



PostgreSQL and GraphRAG Integration in Docker

Presenter: Helen Zeng
AI GBB

May 27, 2025

<https://github.com/Azure-Samples/postgreSQL-graphRAG-docker>



Background and objectives:

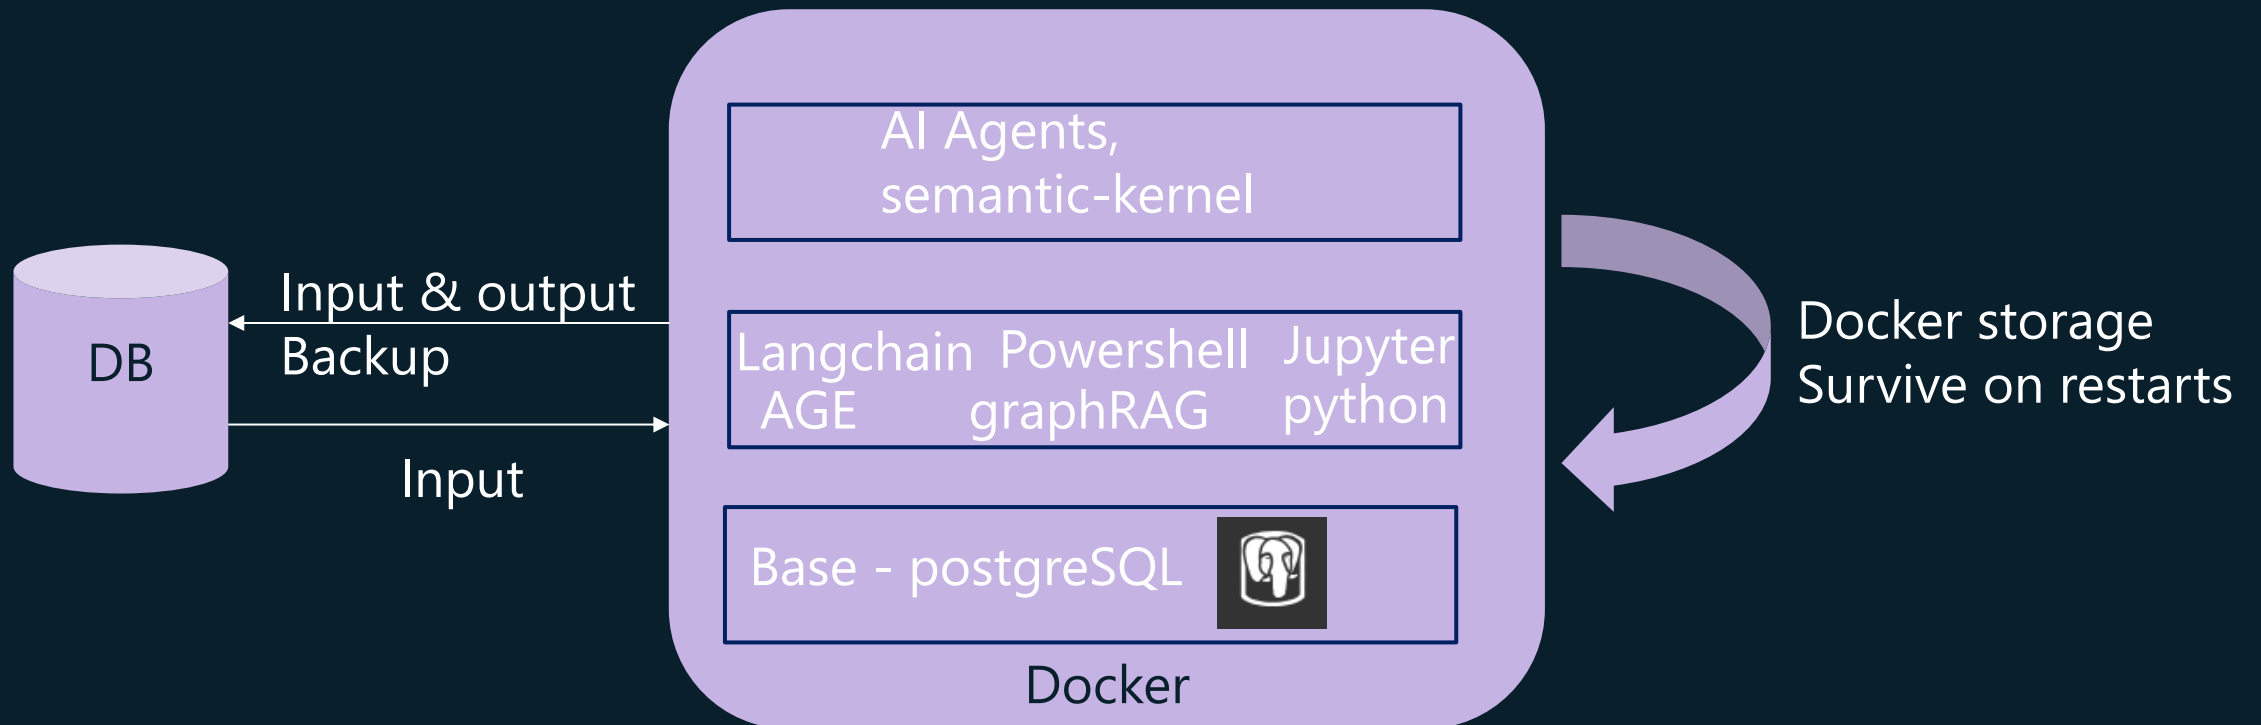
- Raw input data is already in a DB.
- Open Cypher query capability is needed.
- Run CLI version of graphRAG Conveniently.
- No need another data storage between the DB and graphRAG.

Solution design:

- Build graphARG index directly using data in the DB.
- Integrate PostgreSQL and graphRAG in a Docker container, add AGE (Apache Graph Extension), for running Cypher query.
- Bonus: add AI agent services on top.

Architecture

The docker image uses postgres as base, then added python, graphRAG, AGE, semantic-kernel and other needed packages. Two volumes are created to persist postgres data and app related data.



```
volumes:  
  - local_postgres_data:/var/lib/postgresql/data  
  - graphrag_data:/app/graphrag-folder
```

Preparation

1. Prepare Dockerfile
2. Prepare the .env file
3. Prepare setting.yaml
4. Prepare docker-compose.yaml
5. Prepare python code or jupyter notebook

.env example

Your DB properties

Your docker postgres
property

If you use your docker
postgres, then set to true

```
1 MY_DB_HOST=
2 MY_DB_PORT=
3 MY_DB_USER=
4 MY_DB_PASSWORD=
5 MY_DB_NAME=
6 MY_DB_SSLMODE=
7 PROMPTS=/app/graphrag-folder/prompts
8 CONFIG_PATH=/app/graphrag-folder/settings.yaml
9 AGE_HOST=postgres
10 AGE_PORT=5432
11 AGE_USER=postgres
12 AGE_PASSWORD=
13 AGE_DB=postgres
14 USE_LOCAL_AGE=true
15 GRAPHRAG_API_BASE=
16 GRAPHRAG_API_KEY=
17 GRAPHRAG_LLM_MODEL=gpt-4o
18 GRAPHRAG_LLM_API_VERSION="2025-01-01-preview"
19 GRAPHRAG_LLM_DEPLOYMENT=gpt-4o
20 GRAPHRAG_EMBEDDING_MODEL=text-embedding-3-small
21 GRAPHRAG_EMBEDDING_API_VERSION="2023-05-15"
22 GRAPHRAG_EMBEDDING_DEPLOYMENT=text-embedding-3-small
```

Step 0 – insert .txt input to DB if not already existing

Run python code to insert .txt files to DB:

```
> python insert-table.py
```

Step 1 – build docker image

This step builds docker images to include all modules needed to run services later.

```
$ docker build -t graphrag-img .
```

Step 2 – run postgres service

This step starts postgres service which other services depend on.

```
$ docker compose up postgres
```

```
[+] Running 2/2
  ✓ Network postgresql-age_default Created 0.1s
  ✓ Container postgres Created 0.1s
Attaching to postgres
postgres |
postgres | PostgreSQL Database directory appears to contain a database; Skipping initialization
postgres |
postgres | 2025-05-23 16:28:27.226 UTC [1] LOG: starting PostgreSQL 16.9 (Debian 16.9-1.pgdg120+1) on x86_64-pc-linux-gnu, compiled by gcc (Debian 12.2.0-14) 12.2.0, 64-bit
postgres | 2025-05-23 16:28:27.227 UTC [1] LOG: listening on IPv4 address "0.0.0.0", port 5432
postgres | 2025-05-23 16:28:27.227 UTC [1] LOG: listening on IPv6 address ":::", port 5432
postgres | 2025-05-23 16:28:27.234 UTC [1] LOG: listening on Unix socket "/var/run/postgresql/.s.PGSQL.5432"
postgres | 2025-05-23 16:28:27.250 UTC [29] LOG: database system was shut down at 2025-05-23 16:05:41 UTC
postgres | 2025-05-23 16:28:27.270 UTC [1] LOG: database system is ready to accept connections
[]
```


Step 3.1 – run 'load-data' service

This step loads data from DB to Docker folder /app/graphrag-folder/input as graphRAG input.

```
$ docker compose up load-data
```

```
[+] Running 2/2
 ✓ Container postgres      Running      0.0s
 ✓ Container load-data-app Created      0.1s
Attaching to load-data-app
load-data-app | All rows saved as .txt files in /app/graphrag-folder/input.
load-data-app exited with code 0
```

Step 3.2 – check data is present in postgres server

This sub-step checks docker folder /app/graphrag-folder/input to verify graphRAG input.

```
$ docker exec -it postgres bash
```

```
root@b754b0d605cb:/app/graphrag-folder# cd input
root@b754b0d605cb:/app/graphrag-folder/input# ls
input_10.txt input_1.txt input_2.txt input_3.txt input_4.txt input_5.txt input_6.txt input_7.txt input_8.txt input_9.txt
root@b754b0d605cb:/app/graphrag-folder/input#
```

Step 4.1 – build graphrag index

This step runs 'graphrag index' using /app/graphrag-folder/input and generate output.

```
$ docker compose up graphrag-index
```

```
[+] Running 3/3
✓ Container postgres           Running      0.0s
✓ Container load-data-app      Created    0.0s
✓ Container graphrag-index-app Created    0.1s
Attaching to graphrag-index-app
graphrag-index-app | Total files: 10
```

```
graphrag-index-app | [119 rows x 2 columns]]
graphrag-index-app | ** GraphRAG Indexer
graphrag-index-app | — Loading Input (InputFileType.text) - 10 files loaded (0 filtered) - 100% 0...
graphrag-index-app | — create_base_text_units ————— 100% 0:00:00 0:00:00
graphrag-index-app | — create_final_documents ————— 100% 0:00:00 0:00:00
graphrag-index-app | — extract_graph ————— 100% 0:00:00 0:00:00
graphrag-index-app | — finalize_graph ————— 100% 0:00:00 0:00:00
graphrag-index-app | — create_communities ————— 100% 0:00:00 0:00:00
graphrag-index-app | — create_final_text_units ————— 100% 0:00:00 0:00:00
graphrag-index-app | — create_community_reports ————— 100% 0:00:00 0:00:00
graphrag-index-app | — generate_text_embeddings ————— 100% 0:00:00 0:00:00 🚀 All workflows completed successfully.
graphrag-index-app | ccessfully.
graphrag-index-app | graphrag-index-app exited with code 0
```

Step 4.2 – check data is present in postgres server

This sub-step verifies 'graphrag index' output in /app/graphrag-folder/output.

```
!$ docker exec -it postgres bash
```

```
root@b754b0d605cb:/app# cd graphrag-folder
root@b754b0d605cb:/app/graphrag-folder# ls
cache input logs output prompts settings.yaml update_output
root@b754b0d605cb:/app/graphrag-folder# cd input
root@b754b0d605cb:/app/graphrag-folder/input# ls
input_10.txt input_1.txt input_2.txt input_3.txt input_4.txt input_5.txt input_6.txt input_7.txt input_8.txt input_9.txt
root@b754b0d605cb:/app/graphrag-folder/input# cd ../output
root@b754b0d605cb:/app/graphrag-folder/output# ls
communities.parquet      embeddings.community.full_content.parquet  graph.graphml      text_units.parquet
community_reports.parquet embeddings.entity.description.parquet      lancedb
context.json             embeddings.text_unit.text.parquet          relationships.parquet
documents.parquet        entities.parquet                           stats.json
root@b754b0d605cb:/app/graphrag-folder/output# cd lancedb
root@b754b0d605cb:/app/graphrag-folder/output/lancedb# ls
default-community-full_content.lance  default-entity-description.lance  default-text_unit-text.lance
root@b754b0d605cb:/app/graphrag-folder/output/lancedb#
```

Step 5 – write index output to DB

This step stores the graph output data in DB, as docker backup.

The content in /app/graphrag-folder/output will be saved into DB. This example uses Azure hosted PostgreSQL as DB.

```
$ docker compose up graphrag-writer
```

```
[+] Running 4/4
✓ Container postgres      Running      0.0s
✓ Container load-data-app Created      0.0s
✓ Container graphrag-index-app Created      0.0s
✓ Container graphrag-writer-app Created      0.1s
Attaching to graphrag-writer-app
graphrag-writer-app | Skipping JSON file: /app/graphrag-folder/logs/logs.json
graphrag-writer-app | All relevant files have been processed and stored in PostgreSQL.
graphrag-writer-app exited with code 0
```

Step 6 – build graph, prepare for Cypher query

This step builds graph using AGE on PostgreSQL in docker, to prepare for Cypher query.

```
$ docker compose up build-graph
```

```
[+] Running 5/5
✓ Container postgres          Running      0.0s
✓ Container load-data-app     Created      0.0s
✓ Container graphrag-index-app Created      0.0s
✓ Container graphrag-writer-app Created      0.0s
✓ Container build-graph-app   Created      0.1s
Attaching to build-graph-app
build-graph-app | Existing graph 'graphRAG' dropped.
build-graph-app | Graph 'graphRAG' will be created.
build-graph-app | Inserted 10 rows into Document
build-graph-app | Index(['id', 'human_readable_id', 'title', 'text', 'text_unit_ids',
build-graph-app |         'creation_date', 'metadata'],
build-graph-app |         dtype='object')
build-graph-app | Inserted 449 rows into Entity
build-graph-app | Index(['id', 'human_readable_id', 'title', 'type', 'description',
build-graph-app |         'text_unit_ids', 'frequency', 'degree', 'x', 'y'],
build-graph-app |         dtype='object')
build-graph-app | Inserted 690 rows into relationship
build-graph-app | Index(['id', 'human_readable_id', 'source', 'target', 'description', 'weight',
build-graph-app |         'combined_degree', 'text_unit_ids'],
build-graph-app |         dtype='object')
build-graph-app | Data loaded into AGE graph successfully.
build-graph-app exited with code 0
```

View the graph content

This step verifies graph in PostgreSQL ag_catalog in docker, name graphRAG.

```
root@de59c1cc762b:/app# psql -U postgres -d postgres <<EOF > output.txt
> LOAD 'age';
> SET search_path = ag_catalog, "$user", public;
> SELECT * FROM cypher('graphRAG', $$ MATCH (n) RETURN n LIMIT 5 $$) AS (n agtype);
> EOF
```

```
root@de59c1cc762b:/app# less output.txt
```


Step 7 – run query in jupyter notebook

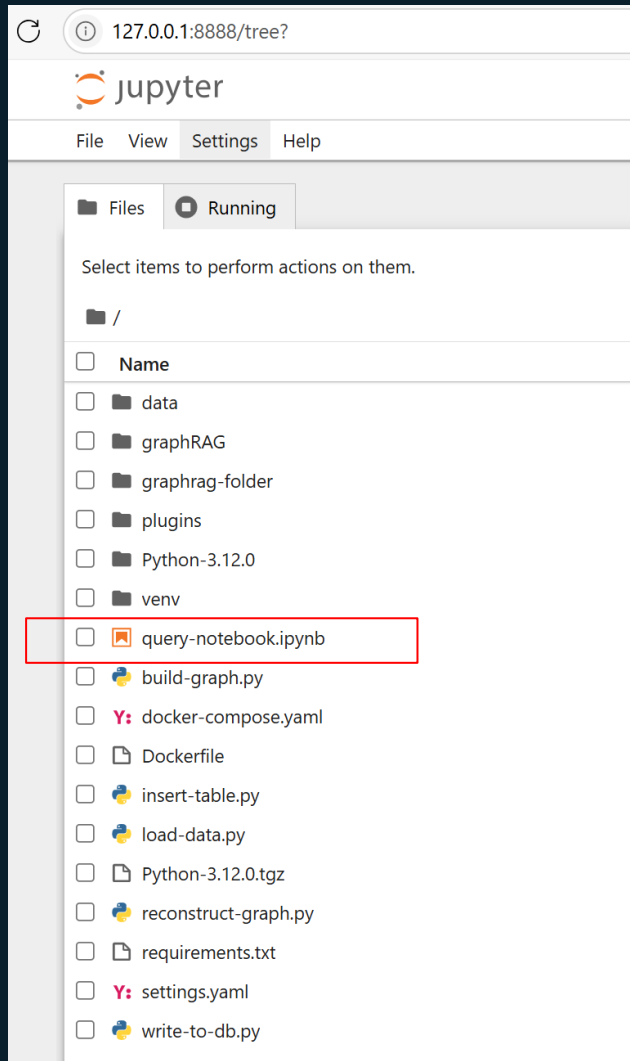
This step runs jupyter notebook in docker.

```
$ docker compose up query-notebook
```

```
[+] Running 6/6
✓ Container postgres          Running          0.0s
✓ Container load-data-app     Created           0.0s
✓ Container graphrag-index-app Created           0.0s
✓ Container graphrag-writer-app Created           0.0s
✓ Container build-graph-app   Created           0.0s
✓ Container query-notebook-app Recreated         0.1s
Attaching to query-notebook-app
query-notebook-app | [I 2025-05-23 16:51:30.721 ServerApp] jupyter_lsp | extension was successfully linked.
query-notebook-app | [I 2025-05-23 16:51:30.727 ServerApp] jupyter_server_terminals | extension was successfully linked.
query-notebook-app | [W 2025-05-23 16:51:30.728 LabApp] 'shutdown_no_activity_timeout' has moved from NotebookApp to ServerApp. This config will be passed to ServerApp. Be sure to update your config before our next release.
query-notebook-app | [I 2025-05-23 16:51:30.730 ServerApp] jupyterlab | extension was successfully linked.
query-notebook-app | [I 2025-05-23 16:51:30.734 ServerApp] notebook | extension was successfully linked.
query-notebook-app | [I 2025-05-23 16:51:30.742 ServerApp] Writing Jupyter server cookie secret to /root/.local/share/jupyter/runtime/jupyter_cookie_secret
query-notebook-app | [I 2025-05-23 16:51:31.018 ServerApp] notebook_shim | extension was successfully linked.
query-notebook-app | [I 2025-05-23 16:51:31.060 ServerApp] notebook_shim | extension was successfully loaded.
query-notebook-app | [I 2025-05-23 16:51:31.062 ServerApp] jupyter_lsp | extension was successfully loaded.
query-notebook-app | [I 2025-05-23 16:51:31.063 ServerApp] jupyter_server_terminals | extension was successfully loaded.
query-notebook-app | [I 2025-05-23 16:51:31.078 LabApp] JupyterLab extension loaded from /app/venv/lib/python3.12/site-packages/jupyterlab
query-notebook-app | [I 2025-05-23 16:51:31.078 LabApp] JupyterLab application directory is /app/venv/share/jupyter/lab
query-notebook-app | [I 2025-05-23 16:51:31.080 LabApp] Extension Manager is 'pypi'.
query-notebook-app | [I 2025-05-23 16:51:31.136 ServerApp] jupyterlab | extension was successfully loaded.
query-notebook-app | [I 2025-05-23 16:51:31.138 ServerApp] notebook | extension was successfully loaded.
query-notebook-app | [I 2025-05-23 16:51:31.139 ServerApp] Serving notebooks from local directory: /app
query-notebook-app | [I 2025-05-23 16:51:31.139 ServerApp] Jupyter Server 2.16.0 is running at:
query-notebook-app | [I 2025-05-23 16:51:31.139 ServerApp] http://
query-notebook-app | [I 2025-05-23 16:51:31.139 ServerApp] http://127.0.0.1:8888/tree?token=
query-notebook-app | [I 2025-05-23 16:51:31.139 ServerApp] Use Control-C to stop this server and shut down all kernels (twice to stop confirmation).
query-notebook-app | [C 2025-05-23 16:51:31.150 ServerApp]
```

Click this link to explore all files in the project folder.

Jupyter notebook UI



The screenshot shows the JupyterLab file browser sidebar. The address bar displays '127.0.0.1:8888/tree?'. The sidebar has tabs for 'Files' and 'Running'. Under 'Files', there is a list of files and folders. The file 'query-notebook.ipynb' is highlighted with a red box. Other files include 'data', 'graphRAG', 'graphrag-folder', 'plugins', 'Python-3.12.0', 'venv', 'build-graph.py', 'docker-compose.yaml', 'Dockerfile', 'insert-table.py', 'load-data.py', 'Python-3.12.0.tgz', 'reconstruct-graph.py', 'requirements.txt', 'settings.yaml', and 'write-to-db.py'.

127.0.0.1:8888/tree?

Jupyter

File View Settings Help

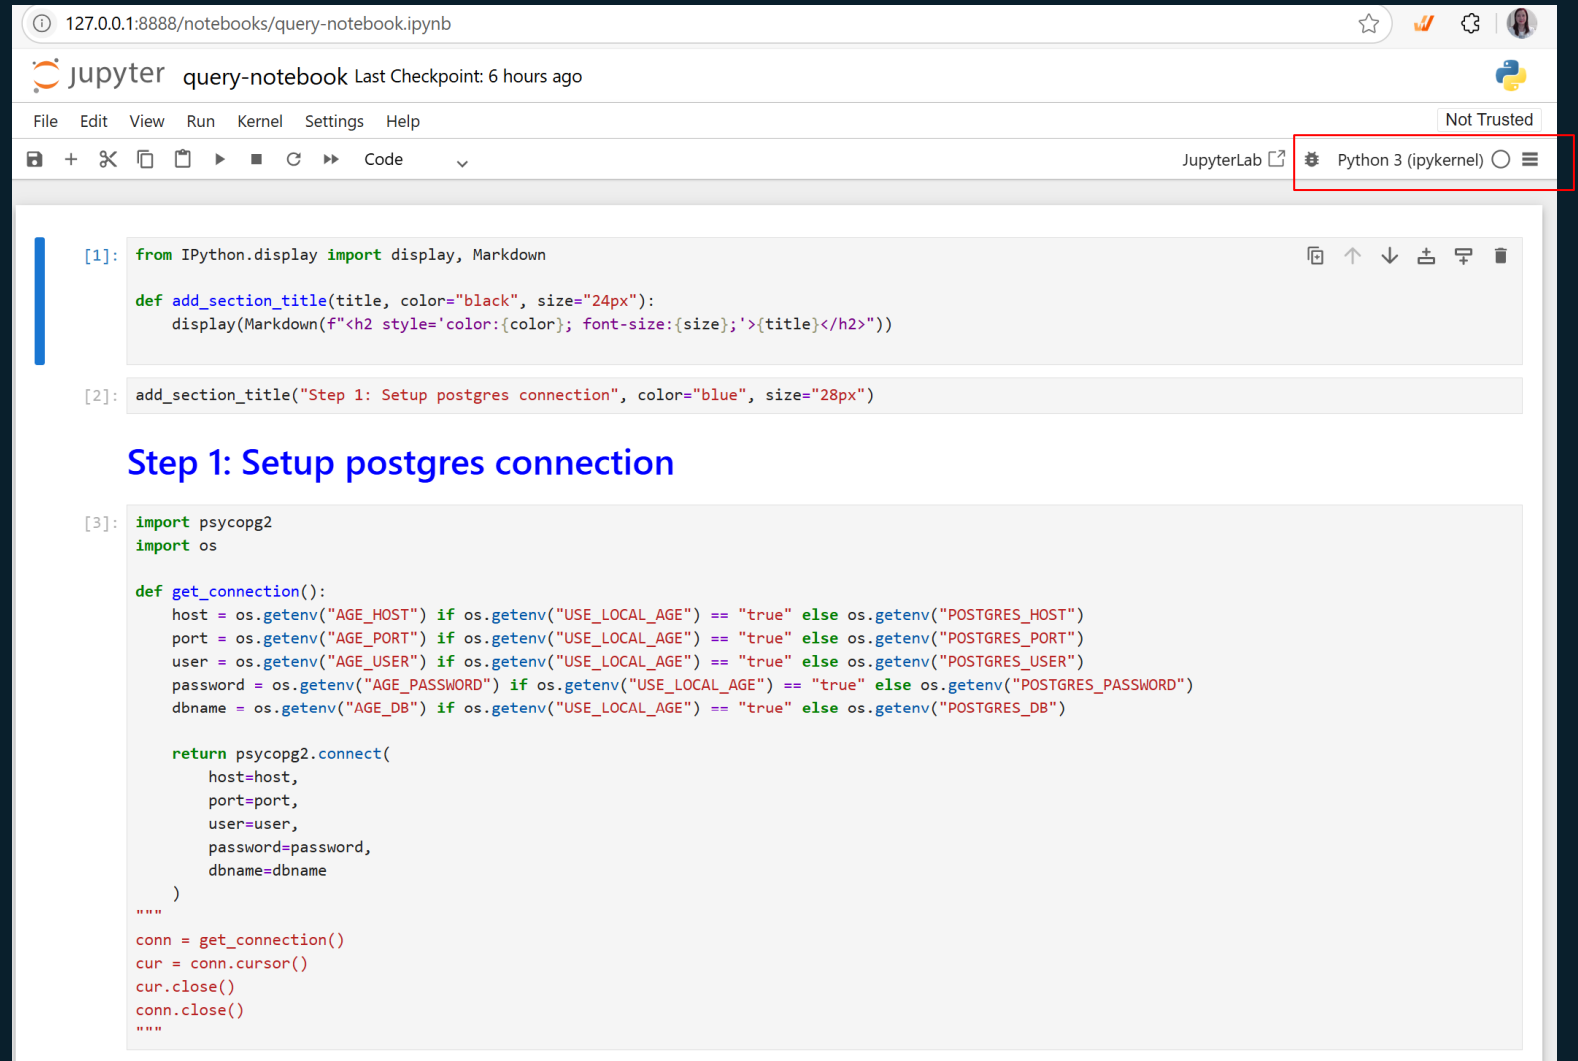
Files Running

Select items to perform actions on them.

/

Name

- ☐ data
- ☐ graphRAG
- ☐ graphrag-folder
- ☐ plugins
- ☐ Python-3.12.0
- ☐ venv
- ☒ query-notebook.ipynb
- ☐ build-graph.py
- ☐ docker-compose.yaml
- ☐ Dockerfile
- ☐ insert-table.py
- ☐ load-data.py
- ☐ Python-3.12.0.tgz
- ☐ reconstruct-graph.py
- ☐ requirements.txt
- ☐ settings.yaml
- ☐ write-to-db.py



The screenshot shows the JupyterLab code editor. The address bar displays '127.0.0.1:8888/notebooks/query-notebook.ipynb'. The editor has tabs for 'File', 'Edit', 'View', 'Run', 'Kernel', 'Settings', and 'Help'. The 'Kernel' tab is active, showing 'Python 3 (ipykernel)' with a red box around it. The code editor contains three cells of code. The first cell defines a function 'add_section_title'. The second cell calls this function. The third cell imports 'psycopg2' and 'os', and defines a function 'get_connection'.

127.0.0.1:8888/notebooks/query-notebook.ipynb

Jupyter query-notebook Last Checkpoint: 6 hours ago

File Edit View Run Kernel Settings Help

Python 3 (ipykernel)

```
[1]: from IPython.display import display, Markdown

def add_section_title(title, color="black", size="24px"):
    display(Markdown(f"<h2 style='color:{color}; font-size:{size};>{title}</h2>"))

[2]: add_section_title("Step 1: Setup postgres connection", color="blue", size="28px")

Step 1: Setup postgres connection

[3]: import psycopg2
import os

def get_connection():
    host = os.getenv("AGE_HOST") if os.getenv("USE_LOCAL_AGE") == "true" else os.getenv("POSTGRES_HOST")
    port = os.getenv("AGE_PORT") if os.getenv("USE_LOCAL_AGE") == "true" else os.getenv("POSTGRES_PORT")
    user = os.getenv("AGE_USER") if os.getenv("USE_LOCAL_AGE") == "true" else os.getenv("POSTGRES_USER")
    password = os.getenv("AGE_PASSWORD") if os.getenv("USE_LOCAL_AGE") == "true" else os.getenv("POSTGRES_PASSWORD")
    dbname = os.getenv("AGE_DB") if os.getenv("USE_LOCAL_AGE") == "true" else os.getenv("POSTGRES_DB")

    return psycopg2.connect(
        host=host,
        port=port,
        user=user,
        password=password,
        dbname=dbname
    )

conn = get_connection()
cur = conn.cursor()
cur.close()
conn.close()
```

Cypher query examples

Several ways of querying are provided:

- Cypher query
 - Example given to check nodes and relationships, and their properties.
- Multi-hop cypher query
 - One of the advantages of using knowledge graph is the multi-hop query.
- GraphRAG LocalSearch
 - Examples given to show the impact of community level to the query results.
- GraphRAG GlobalSearch
 - Both python function call and CLI methods are provided.
- Vector search
 - Example shows that vector search result is restricted to the original text; it typically returns the top-k most similar original chunks. In other words, vector search is retrieval only, it returns what was said, not what it means. However, graphRAG combines vector search with a knowledge graph that encodes relationships between entities. The graph is built from LLM-extracted triples, which helps contextual thinking; thus, graphRAG provides more synthesized and thoughtful response.

Build and run AI agents example

Task: summarize the podcasts.

```
await summarize_all_text("/app/graphrag-folder/input", agent)
```

--- Combined Summary ---

Sure, here's a concise summary of the lengthy text you provided:

****Summary:****

The podcast **Behind the Tech**, hosted by Microsoft CTO Kevin Scott, explores diverse stories about innovators in technology. In this series, Scott interviews influential figures such as Jaron Lanier, Reid Hoffman, Danielle Feinberg, Alice Steinglass, Andrew Ng, Anders Hejlsberg, Surya Ganguli, and others, showcasing their unique journeys into tech and their impact on the industry.

Common themes across episodes include:

1. ****Career Journeys:**** Interviewees shared personal anecdotes about how curiosity, chance, and determination shaped their paths—from early exposure to computing to tackling complex challenges in their careers, including machine learning, VR, programming languages, and animation.
2. ****Intersection of Creativity and STEM:**** Guests like Danielle Feinberg and Anders Hejlsberg describe blending artistic and technical craftsmanship in their work, whether creating visually stunning Pixar films or designing programming tools like Turbo Pascal and C#.
3. ****Advocacy for Diversity & Education:**** Many guests emphasize the importance of inspiring younger generations, especially women and minorities, to pursue computer science. Initiatives like **Code.org** and Stanford's Human-Centered AI Institute are addressing systemic gaps to improve diversity and accessibility in tech.
4. ****Ethical & Societal Implications:**** Episodes discuss AI's rapid evolution, needing thoughtful regulation and ethical responsibility. Insights from academics like Surya Ganguli highlight concerns ranging from biases in data to the energy inefficiencies of digital computation.
5. ****Technological Impact:**** Guests explore transformative industries such as VR, healthcare, fintech, and agriculture, demonstrating ways technology solves real-world problems and shapes global economies.

Interview highlights include:

- Jaron Lanier's critique of the Internet's influence on society.
- Reid Hoffman's insights on entrepreneurial success and China's ability to "blitzscale."
- Danielle Feinberg's role in leveraging math and storytelling in Pixar animation.
- Alice Steinglass advocating CS education for diverse classrooms.

Through their candid experiences, these technologists inspire listeners to appreciate the challenges, opportunities, and ever-expanding possibilities within STEM.

Let me know if you'd like to dive deeper into any individual story from the text.

Optional step: reconstruct-graph

In case the graph related data in docker /app/graphrag-folder/output lost, then the graph can be reconstructed using the data stored in DB, which was saved in step 5.

```
$ docker compose up reconstruct-graph
```

```
[+] Running 2/2
 ✓ Container postgres          Running      0.0s
 ✓ Container reconstruct-graph-app Created    0.3s
Attaching to reconstruct-graph-app
reconstruct-graph-app | All files have been reconstructed from Azure-hosted PostgreSQL to /app/graphrag-folder/re
store.
reconstruct-graph-app exited with code 0
```

```
$ docker exec -it postgres bash
```

```
root@076170b43c6f:/app# cd graphrag-folder
root@076170b43c6f:/app/graphrag-folder# ls
cache graph.graphml input logs output plugins prompts restore settings.yaml update_output
root@076170b43c6f:/app/graphrag-folder# cd restore
root@076170b43c6f:/app/graphrag-folder/restore# ls
communities.parquet      embeddings.community.full_content.parquet  graph.graphml      text_units.parquet
community_reports.parquet embeddings.entity.description.parquet      relationships.parquet
context.json            embeddings.text_unit.text.parquet          stats.json
documents.parquet       entities.parquet                          summaries.json
root@076170b43c6f:/app/graphrag-folder/restore#
```

Query in docker postgres

```
$ docker exec -it postgres bash
```

```
root@307e79cfb708:/app# psql -U postgres -d postgres
psql (16.9 (Debian 16.9-1.pgdg120+1))
Type "help" for help.
```

```
postgres=# LOAD 'age';
```

```
LOAD
```

```
postgres=# SET search_path = ag_catalog, "\user", public;
SET
```

```
postgres=# SELECT * FROM cypher('graphRAG', $$
    MATCH (n) RETURN labels(n), count(n)
    $$) AS (labels agtype, count agtype);
 labels | count
```

```
-----+-----
["Entity"] | 449
["Document"] | 10
(2 rows)
```

```
postgres=# SELECT * FROM cypher('graphRAG', $$
    MATCH ()-[r]->() RETURN type(r), count(r)
    $$) AS (relationship agtype, count agtype);
 relationship | count
```

```
-----+-----
"RELATED_TO" | 690
(1 row)
```

```
postgres=# \dt
```

```
      List of relations
 Schema | Name   | Type  | Owner
-----+-----+-----+-----
 ag_catalog | ag_graph | table | postgres
 ag_catalog | ag_label | table | postgres
(2 rows)
```

```
postgres=# █
```

Graph storage

Graph metadata, like nodes, edges stored in Docker postgres

```
$ docker exec -it postgres bash
```

```
root@299fc1439d4d:/app# psql -U postgres -d postgres
psql (16.9 (Debian 16.9-1.pgdg120+1))
Type "help" for help.

postgres=# \dt
Did not find any relations.
postgres=# LOAD 'age';
LOAD
postgres=# SET search_path = ag_catalog, "$user", public;
SET
postgres=# SELECT * FROM cypher('graphRAG', $$
    MATCH ()--() RETURN count(*)
$$) AS (count agtype);
count
-----
1380
(1 row)

postgres=# SELECT * FROM cypher('graphRAG', $$
    MATCH ()-[r]->() RETURN count(r)
$$) AS (count agtype);
count
-----
690
(1 row)
```

Graph input and output stored in DB (in this example, it's Azure hosted PostgreSQL)

```
postgres=> \dt
List of relations
Schema | Name | Type | Owner
-----+-----+-----+-----
public | graphrag_inputs | table | XXXXX
public | graphrag_outputs | table | XXXX
(2 rows)
```